

future.mirai: Use the Mirai Parallelization Framework in Futureverse - Easy!

Henrik Bengtsson

University of California, San Francisco

R Foundation

R Consortium

 @HenrikBengtsson

useR! 2024, Salzburg, Austria (2024-07-09)



Futureverse - A Friendly, Unifying Parallelization Framework in R²

- Package **future** provides fundamental building blocks for evaluating R code in parallel
 - **future()**, **value()**, and **resolved()**
 - **%<-%** (future-assignment operator) on top of **future()** & **value()**

```
1 > x <- 1:100
2 > y <- slow_sum(x)          # ~1 min ... waiting!
3 > y
4 [1] 5050
```



Total time: 1 minute

Futureverse - A Friendly, Unifying Parallelization Framework in R

- Package **future** provides fundamental building blocks for evaluating R code in parallel
 - **future()**, **value()**, and **resolved()**
 - **%<-%** (future-assignment operator) on top of **future()** & **value()**

```
1 > x <- 1:100
2 > a <- slow_sum(x[ 1:50 ] )      # ~30 sec
3 > b <- slow_sum(x[51:100])       # ~30 sec
4 > y <- a + b
5 > y
6 [1] 5050
```



Total time: 1 minute

Evaluate R expressions in the background

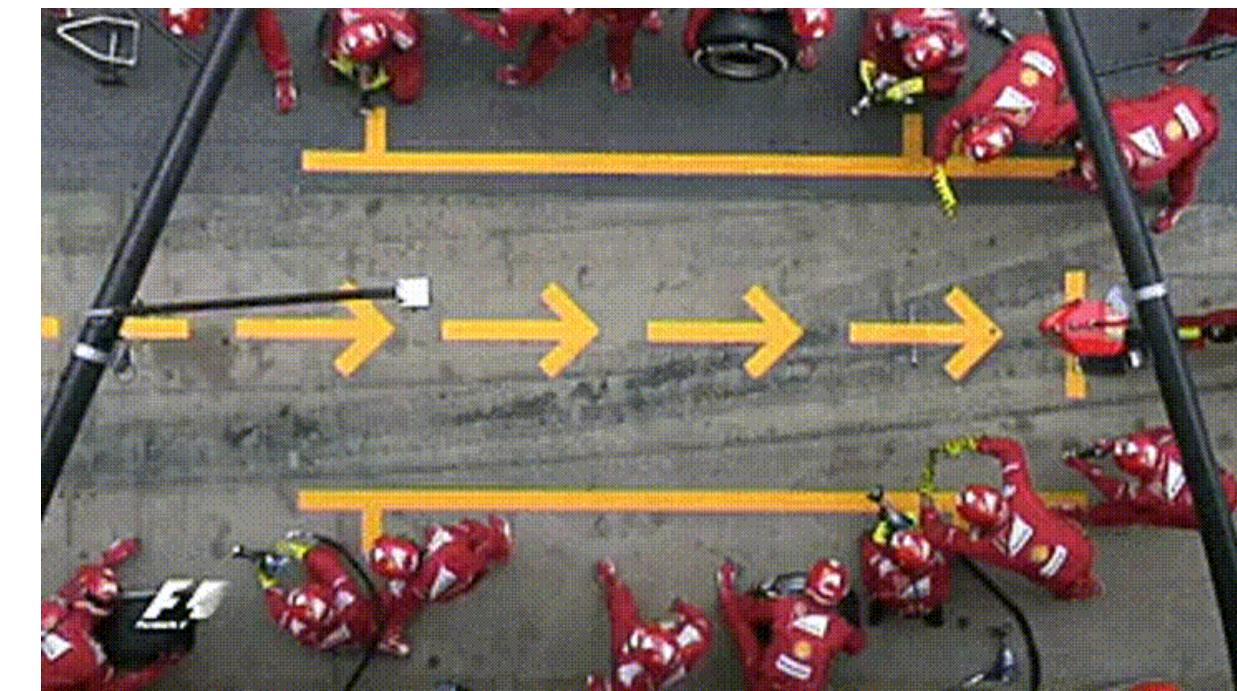
```
1 > library(future)
2 > plan(multisession)                      # parallel
3
4 > x <- 1:100
5 > a %<-% slow_sum(x[ 1:50 ] )          # ~0 sec
6 > b %<-% slow_sum(x[51:100])           # ~0 sec
7
8 > 1 + 2
9 [1] 3
10
11 > y <- a + b                          # get result
12 > y
13 [1] 5050
```



Total time: 30 seconds

Splitting up into more chunks to speed it up further

```
1 > library(future)
2 > plan(multisession)
3
4 > x <- 1:100
5 > a %<-% slow_sum(x[ 1:25 ]) # ~0 sec
6 > b %<-% slow_sum(x[26:50 ]) # ~0 sec
7 > c %<-% slow_sum(x[51:75 ]) # ~0 sec
8 > d %<-% slow_sum(x[76:100]) # ~0 sec
9
10 > y <- a + b + c + d           # get result
11 > y
12 [1] 5050
```



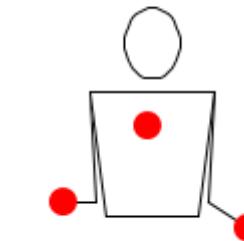
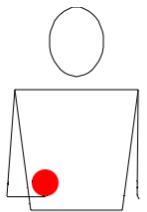
Total time: 15 seconds

End-user can choose from many parallel backends

```
1 plan(sequential)
2 plan(multisession)          # uses {parallel}'s "snow" machinery
3 plan(multicore)            # uses {parallel}'s "multicore" machinery
4
5 plan(cluster, workers = c("n1", "n1", "n1", "n2", "n3"))
6 plan(cluster, workers = c("n1", "m2.uni.edu", "vm.cloud.org"))
```

These are internally based on the **parallel** package.

Higher-level parallelization from `future()` and `value()`



```
1 y <- lapply(X, slow_sum)
```

```
1 plan(multisession, workers = 4)
2 y <- future_lapply(X, slow_sum)
```

Easily implemented via `future()` and `value()`, e.g.

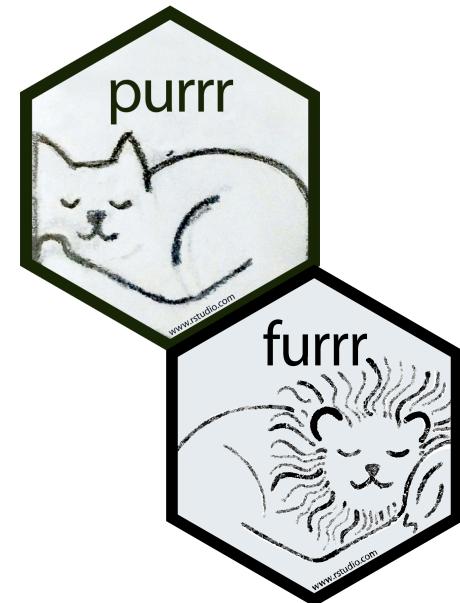
```
1 future_lapply <- function(X, FUN, ...) {
2   fs <- lapply(X, function(x) future(FUN(x, ...)))
3   lapply(fs, value)
4 }
```

User-friendly, higher-level functions

- The concept of “futures” was invented in 1975 (sic!)
- `future()`, `value()`, and `resolved()` are easy to understand, powerful constructs

These building blocks lay the foundation for higher-level functions:

- `future.apply`, e.g. `future_lapply()` and `future_replicate()`
- `furrr`, e.g. `future_map()` and `future_map_dbl()`
- `doFuture`, e.g. `foreach() %dofuture% { ... }`
- ...
- - *Maybe your package will be next?*



Futureverse allows you to stick with your favorite coding style

Parallel alternatives to traditional, sequential functions:

```
1 y <- lapply(x, some_fcn)                                ## base R
2 y <- future_lapply(x, some_fcn)                            ## {future.apply}
```

```
1 y <- map(x, some_fcn)                                     ## {purrr}
2 y <- future_map(x, some_fcn)                               ## {furrr}
```

```
1 y <- foreach(z = x) %do% some_fcn(x)                      ## {foreach}
2 y <- foreach(z = x) %dofuture% some_fcn(z)                ## {foreach} + {doFuture}
```

Yes, we can of course use base-R or **magrittr** pipes where we want to, e.g.

```
1 y <- x |> future_map(some_fcn)
```

```
1 y <- x %>% future_map(some_fcn)
```

Anyone can develop additional parallel backends

From the very beginning in 2015, the plan and hope has been that additional R backends would become available in the future (pun intended!)

The **future.callr** package wraps the **callr** package that is an alternative to the built-in **parallel**-based **multisession** backend:

```
1 plan(future.callr::callr) # locally using callr
```

The **future.batchtools** package wraps the **batchtools** package that can run tasks on high-performance compute (HPC) clusters, e.g.

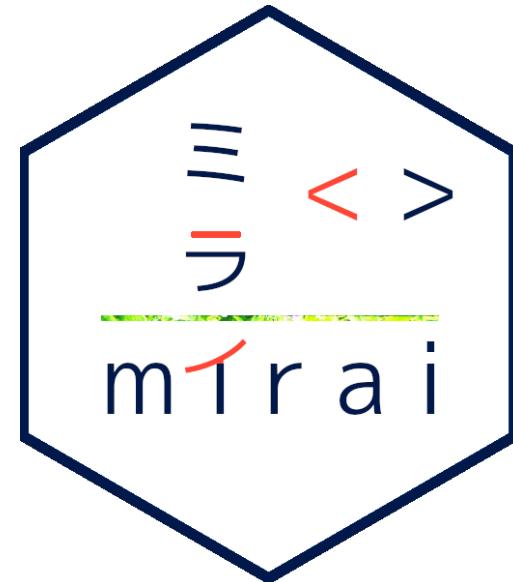
```
1 plan(future.batchtools::batchtools_slurm) # on a Slurm job scheduler  
2 plan(future.batchtools::batchtools_sge) # on a SGE job scheduler
```

And, now also **mirai**-based backends:

```
1 plan(future.mirai::mirai_multisession) # locally using mirai  
2 plan(future.mirai::mirai_cluster) # using mirai cluster
```

mirai - Minimalist Async Evaluation Framework for R

- The **mirai** R package by Charlie Gao (anno 2022)
- *mirai* is Japanese for “future”



Mirai API

```
m <- mirai(expr)
```

create a future

Future API

```
r <- !unresolved(m)
```

check if done

```
f <- future(expr)
```

```
v <- m[]
```

wait & get result

```
r <- resolved(f)
```

```
v <- value(f)
```

- Somewhat lower-level interface than the **future** package
- Minimum overhead through highly optimized implementation
- Provides a powerful queueing-mechanism for processing tasks in parallel

future.mirai - A mirai-based parallel backend for Futureverse

- Makes mirai ecosystem available to Futureverse
- Existing Futureverse code can use it without modification

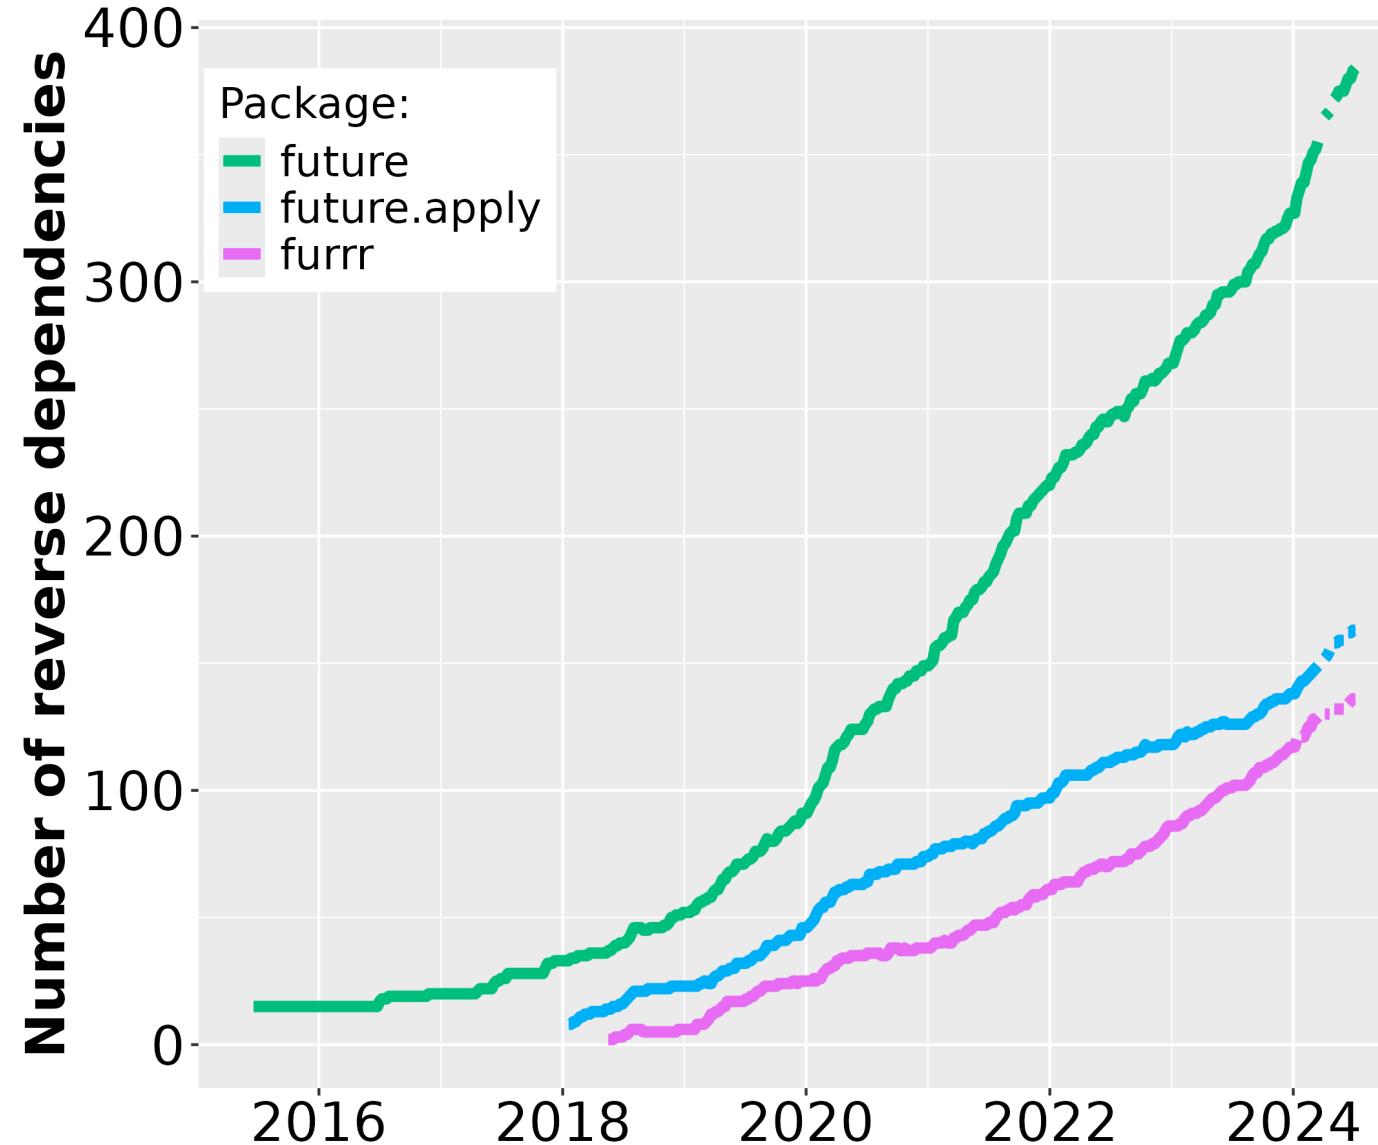
```

1 library(future.mirai)
2 plan(mirai_multisession)      # parallelize via
3                                         # mirai framework
4 x <- rnorm(100)
5 a %<-% sum(x[ 1:50 ])
6 b %<-% sum(x[51:100])
7 y <- a + b
8
9 z <- future.apply::future_sapply(x, slow)
10
11 z <- x |> furrr::future_map_dbl(slow)
12
13 z <- foreach::foreach(.x = x) %dofuture% { slo

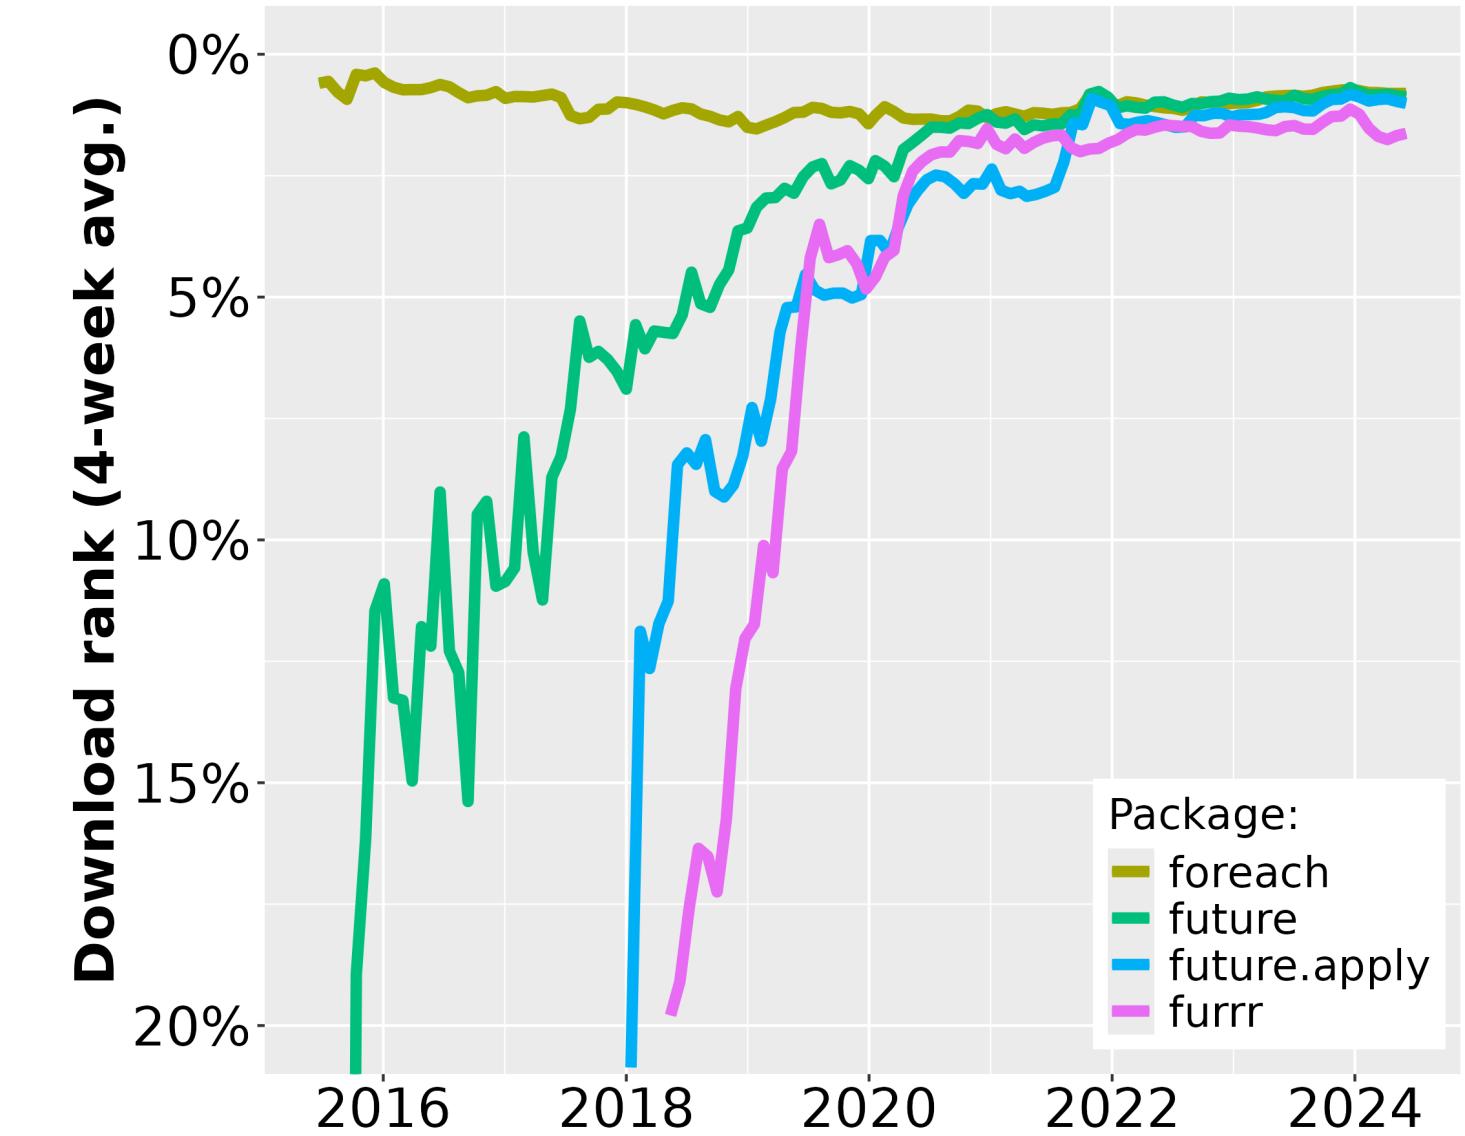
```



Futureverse is very well tested thanks to lots of real-world use



~400 CRAN packages depend directly on the **future** package - it grows 3x faster than **foreach** at 1200 reverse dependencies



The **future** packages is among the 1% most downloaded CRAN packages

As a Futureverse user, you can help mirai!

If you use one of the 100's of CRAN packages that parallelize via Futureverse, by setting:

```
plan(future.mirai::mirai_multisession)
```

you will parallelize via **mirai**, with all its benefits.

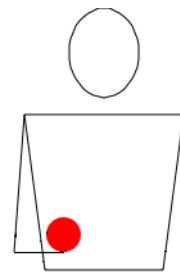
Importantly, by doing so, you will also:

- increase the real-world test coverage of **mirai**
- help increase the stability and quality of **mirai**

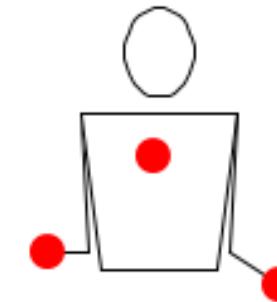
Please give feedback and reach out if you run into issues 

Nested parallelization with some care

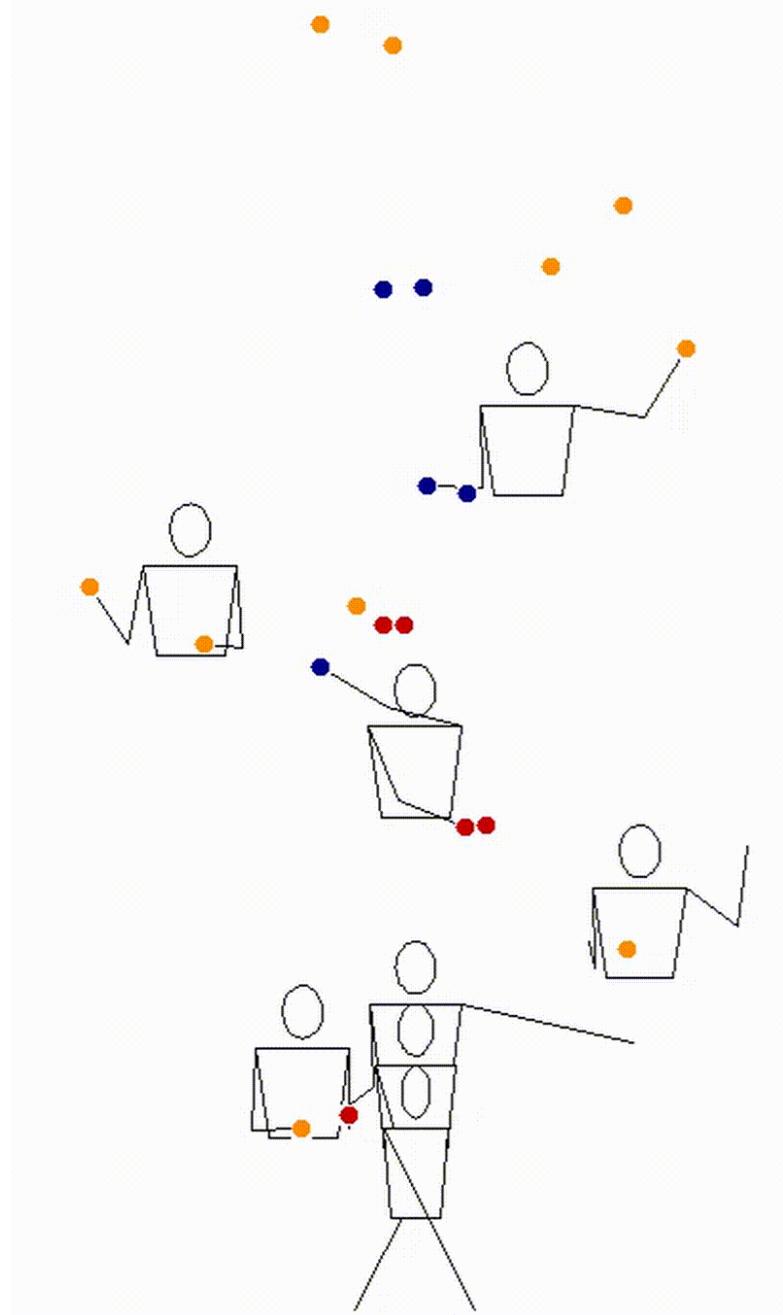
- Futureverse protects against over parallelization
- **future.mirai** opens up for more nested parallelization



Sequential
processing



Parallelization
with 4 workers



Nested parallelization
with 5 workers each
running 3 parallel tasks

Not a competition: Should I use Future API or Mirai API?

- Futureverse is well established and well tested
 - `future.apply`, `furrr`, `foreach` with `doFuture`, ...
 - automatically relays output, messages, warnings, errors, etc.
 - real-time progress updates
 - 100's of packages already parallelize via futures
⇒ they can all use `plan(mirai_multisession)` immediately
- `mirai` is self-contained implementation
 - optimized for minimum overhead
 - undergoes stunning development



There is a promising future for parallelization in R

- It's easy to get started - just try it
- Support: <https://github.com/HenrikBengtsson/future/discussions>
- Tutorials: <https://www.futureverse.org/tutorials.html>
- Blog posts: <https://www.futureverse.org/blog.html>
- More features on the roadmap
- I love feedback and ideas

